

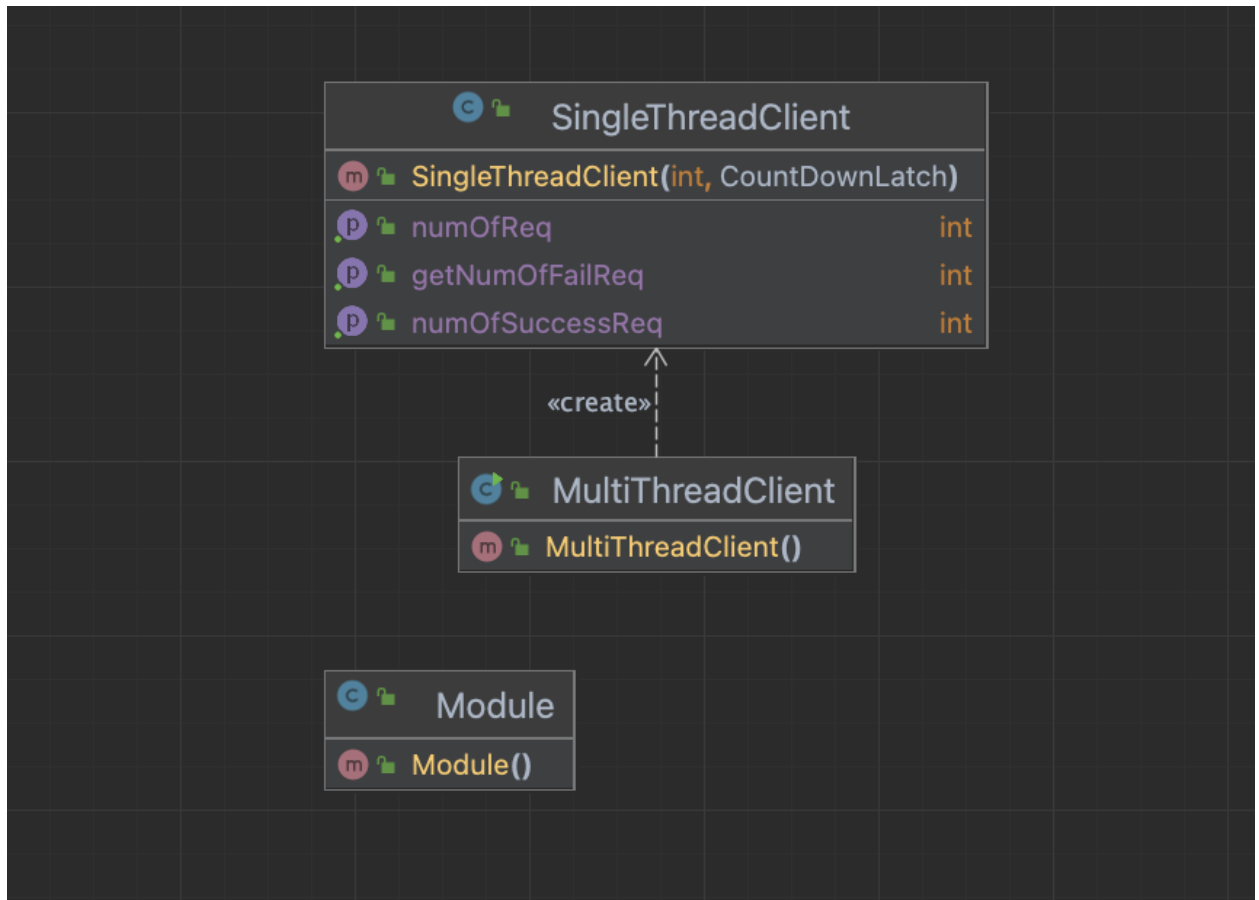
# Submission Requirements

Submit your work to Canvas Assignment 1 as a pdf document. The document should contain:

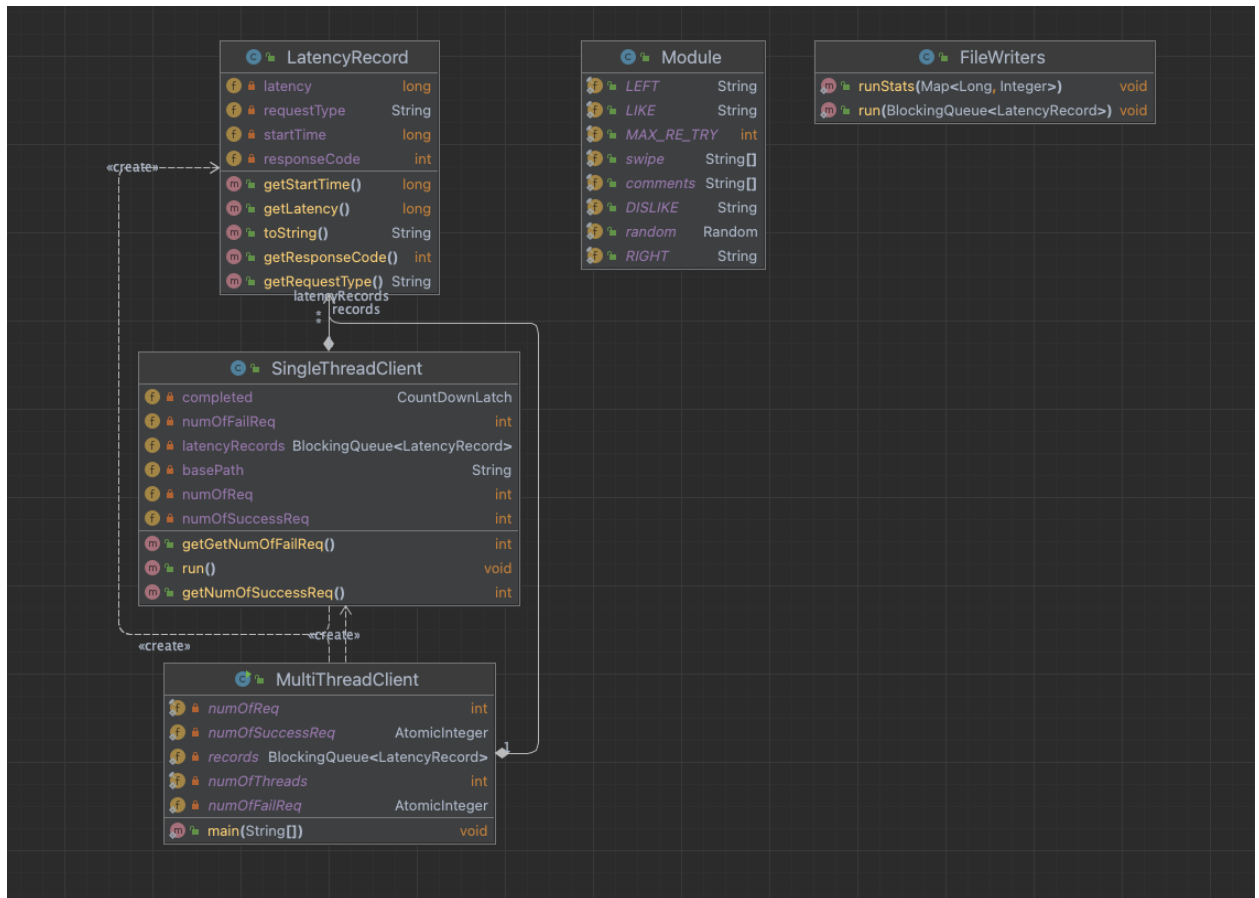
1. the URL for your git repo. *Make sure that the code for the client part 1 and part 2 are in seperate folders in your repo*
2. a 1-2 page description of your client design. Include major classes, packages, relationships, whatever you need to convey concisely how your client works.
3. Client (Part 1) - This should be a screen shot of your output window with your wall time and throughput. Also make sure you include the client configuration in terms of number of threads used and Little's Law throughput predictions.
4. Client (Part 2) - run the client as per Part 1, showing the output window for each run with the specified performance statistics listed at the end.
5. The plot of your throughput over time.

## Submission details:

1. Git repo URL: <https://github.com/Oliver1024/DistributedSystemDatingApp>
2. Client design description:  
Part -1 uml:



Part -2 uml:



The client part which I have is implemented multi-threading client-server architecture to simulate the client sending requests to the server. This client part2 contains 5 Java classes: FileWriters, LatencyRecord, Module, and MultiThreadClient and SingleThreadClient. I also have the same logic as in client part1.

**SingleThreadClient** class acts as a single client, which is executed by multiple threads. Each thread executes the run method of SingleThreadClient class and sends the number of requests specified in the numOfReq variable and the run method of the SingleThreadClient class sends the request to the server, records the latency of each request, and calculates the success and failure counts of the requests. This class creates an instance of a SwipeApi client using the Swagger API client library and sends a specified number of "swipe" requests to the API at the specified base URL. Each swipe request is represented by a "SwipeDetails" object, which includes the swiper ID, swipee ID, and a comment, all of which are randomly generated. If a swipe request fails, the code retries it a specified number of times before moving on to the next request.

**FileWriters** is a class that writes data to two separate CSV files. The first file records latency information, including "StartTime", "RequestType", "Latency", and "ResponseCode". The second file records performance statistics, including "Seconds" and "Throughput/second".

**LatencyRecord** is a simple data class that holds information about a latency record, including the start time, request type, latency, and response code.

**Module** is a class with several constants and a random number generator. The constants define swipe directions (LEFT, RIGHT), comments (LIKE, DISLIKE), and the maximum number of retries (MAX\_RE\_TRY).

**MultiThreadClient** is the main class that creates and runs multiple SingleThreadClient threads. Each SingleThreadClient makes multiple requests and records latency information and performance statistics. The MultiThreadClient class waits for all SingleThreadClient threads to complete and calculates the overall success rate of the requests and the wall-clock time. The final results are stored in two CSV files by the FileWriters class.

### 3. Client - Part 1:

Configure for client:

Number of threads: 100

Number of request per thread: 5000

Total requests: 500k

Thoroughput: 3311 req/s

The screenshot shows an IDE with a project structure on the left and a code editor on the right. The project structure includes a 'java' package with 'org.example' and 'Part1' sub-packages. 'Part1' contains 'Module', 'MultiThreadClient', and 'SingleThreadClient'. The 'MultiThreadClient' class is selected, and its code is displayed in the editor. The code defines constants for the number of threads (100) and requests (5000), initializes atomic counters for successful and failed requests, and implements a main method that starts 100 threads to send requests. The execution output at the bottom shows the results of running the program.

```
5  
6 public class MultiThreadClient {  
7     private static final int numOfThreads = 100;  
8     private static final int numOfReq = 5000;  
9     private static AtomicInteger numOfSuccessReq = new AtomicInteger();  
10    private static AtomicInteger numOfFailReq = new AtomicInteger();  
11  
12    public static void main(String[] args) throws InterruptedException {  
13        long start = System.currentTimeMillis();  
14        CountDownLatch completed = new CountDownLatch(numOfThreads);  
15  
16        SingleThreadClient[] singleThreadClients = new SingleThreadClient[numOfThreads];  
17  
18        for(int i = 0; i < numOfThreads; i++) {  
19            SingleThreadClient singleThreadClient = new SingleThreadClient(numOfReq, completed);  
20            Thread thread = new Thread(singleThreadClient);  
21            singleThreadClients[i] = singleThreadClient;  
22            thread.start();  
23        }  
24  
25        completed.await(); // wait all threads completed  
26    }  
27 }
```

Execution Output:

```
Number of successful requests: 500000  
Number of fail requests: 0  
walltime: 151 seconds  
Total throughput: 3311 req/s  
Process finished with exit code 0
```

First, I send 10000 requests using 1 thread, and get the response time(w) = 18.6ms

Then, apply to little's law,

Expected Little's Law throughput =  $100/18.6 = 5376$  req/s

If I used a consumer-producer pattern, it could help me get better results and be closed to the little's laws' expected results. The producer generates threads first and then adds them to a blocking queue, and the consumer retrieves threads from the blocking queue and processes them.

#### 4. Client - Part2: 100 threads

```
Total throughput: 3333 req/s
Mean response time: 29 ms
Median response time: 28 ms
99th percentile response time: 76 ms
Max response time: 1882 ms
Min response time: 12 ms

Process finished with exit code 0
```

I have conducted numerous experiments to send 500,000 requests, utilizing various configurations such as 10, 20, 30, 50, 80, 100, 200, 250, and 300 threads. Please refer to the following detailed comparison for more information.

numOfThreads: 10  
Total throughput: 605 req/s

```
Total throughput: 605 req/s
Mean response time: 16 ms
Median response time: 15 ms
99th percentile response time: 29 ms
Max response time: 1226 ms
Min response time: 11 ms

Process finished with exit code 0
```

numOfThreads: 20  
Total throughput: 1196 req/s

```
Total throughput: 1196 req/s
Mean response time: 16 ms
Median response time: 16 ms
99th percentile response time: 31 ms
Max response time: 672 ms
Min response time: 11 ms

Process finished with exit code 0
```

numOfThreads: 30

Total throughput: 1724 req/s

```
Total throughput: 1724 req/s
Mean response time: 17 ms
Median response time: 16 ms
99th percentile response time: 35 ms|
Max response time: 839 ms
Min response time: 11 ms

Process finished with exit code 0
```

numOfThreads: 50

Total throughput: 2538 req/s

```
Total throughput: 2538 req/s
Mean response time: 19 ms
Median response time: 18 ms
99th percentile response time: 46 ms
Max response time: 1309 ms
Min response time: 12 ms

Process finished with exit code 0
```

numOfThreads: 80

Total throughput: 3184 req/s

```
Total throughput: 3184 req/s
Mean response time: 24 ms
Median response time: 23 ms
99th percentile response time: 70 ms
Max response time: 1401 ms
Min response time: 12 ms

Process finished with exit code 0
```

numOfThreads: 100

Total throughput: 3333 req/s

```
Total throughput: 3333 req/s
Mean response time: 29 ms
Median response time: 28 ms
99th percentile response time: 76 ms
Max response time: 1882 ms
Min response time: 12 ms

Process finished with exit code 0
```



numOfThreads: 200

Total throughput: 3225 req/s

```
Part2.MultiThreadClient x
/Library/Java/JavaVirtualMachines/jdk-11.0.11.jdk/Contents/Home/bin/java ...
Total throughput: 3225 req/s
Mean response time: 46 ms
Median response time: 42 ms
99th percentile response time: 126 ms
Max response time: 3250 ms
Min response time: 12 ms

Process finished with exit code 0
```

numOfThreads: 250

Total throughput: 3246 req/s

```
Total throughput: 3246 req/s
Mean response time: 76 ms
Median response time: 69 ms
99th percentile response time: 186 ms
Max response time: 5153 ms
Min response time: 13 ms

Process finished with exit code 0
```

numOfThreads: 300

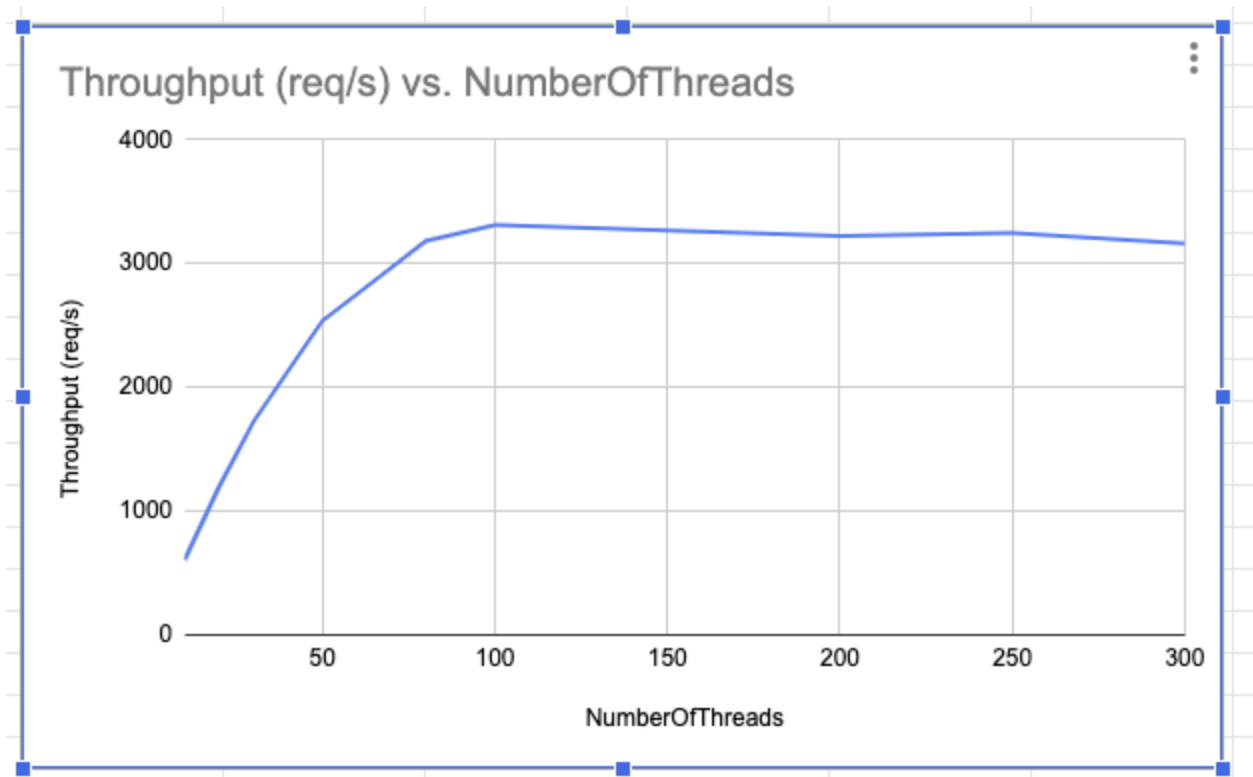
Total throughput: 3164 req/s

```
Total throughput: 3164 req/s
Mean response time: 92 ms
Median response time: 83 ms
99th percentile response time: 235 ms
Max response time: 5288 ms
Min response time: 12 ms

Process finished with exit code 0
```

Finally, I obtained intriguing statistics regarding the relationship between the number of threads and throughput. Surprisingly, the throughput significantly increased when I increased the number of threads from 10 to 100. However, despite the default maximum number of threads for Tomcat being set at 200, the best throughput was achieved with 100 threads for my configuration and client. When I tried to increase the number of threads from 100 to 300, the throughput did not increase and instead, it even decreased slightly.

Additionally, other factors such as AWS networking and time of day may also impact the throughput. For example, the results may vary if the test is conducted in the early morning or late night.



Task 4:

