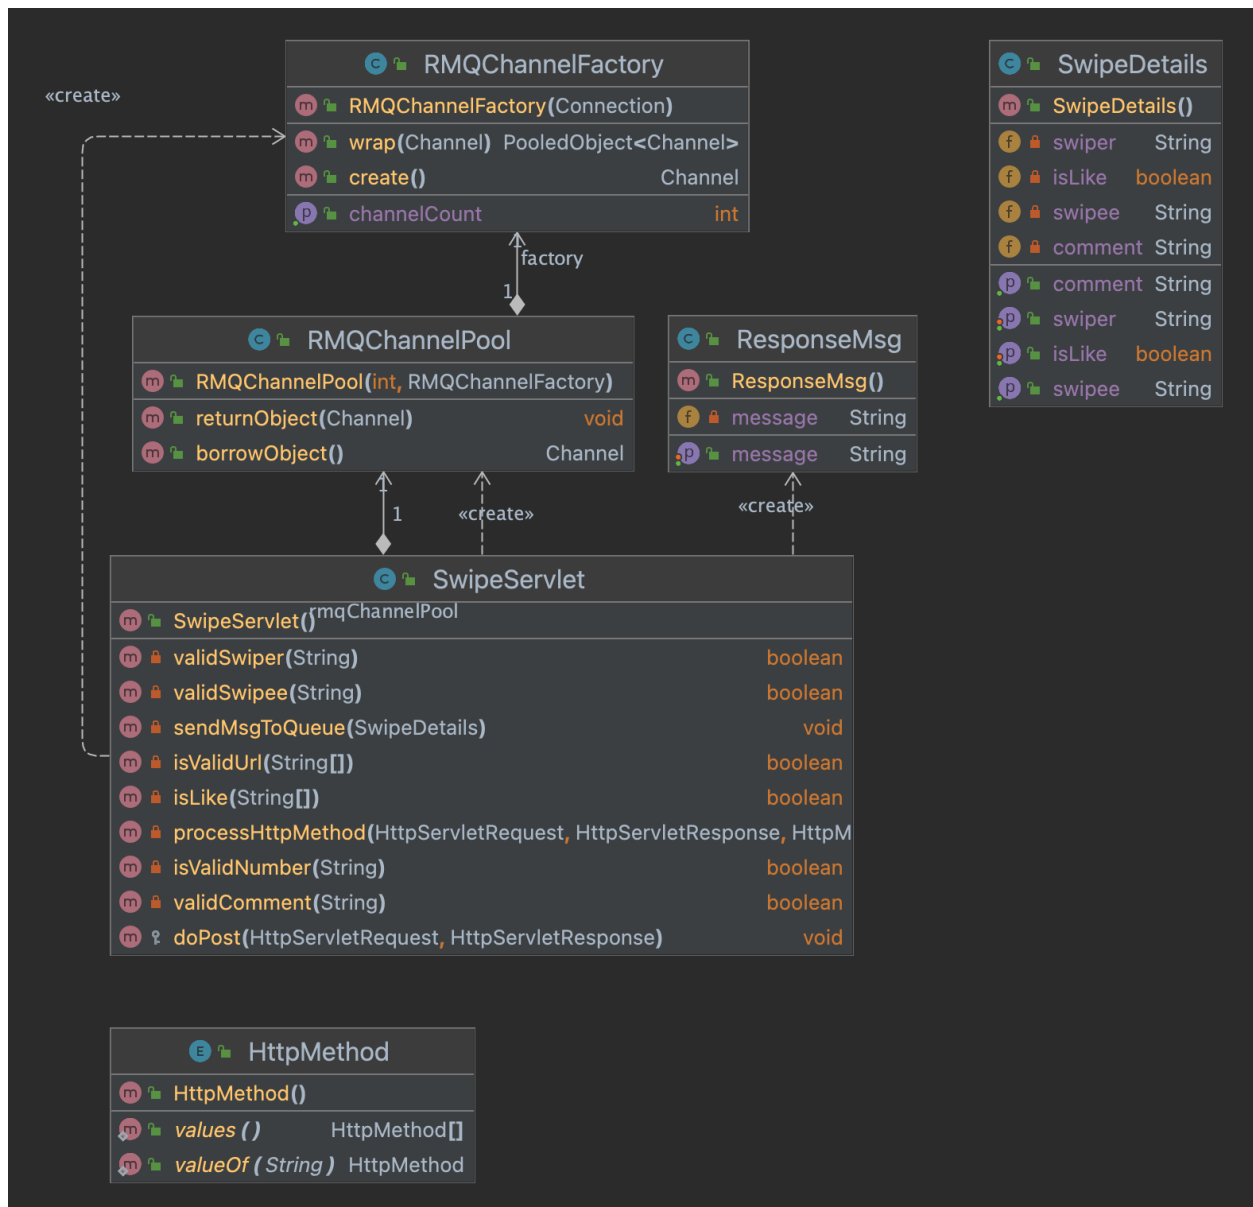


1. Assignment 2 repo

<https://github.com/Oliver1024/DistributedSystemDatingApp/tree/main/assignment2>

2. A 1-2 page description of your server design. Include major classes, packages, relationships, how messages get sent/received, etc



RMQChannelFactory:

The RMQChannelFactory class is a factory for creating Channel objects that are used to communicate with a RabbitMQ server, which implements the BasePooledObjectFactory interface provided by Apache Commons Pool 2, which defines a set of methods for creating, validating, and destroying objects that are managed by an object pool. And the the RMQChannelFactory class has a single constructor that takes a Connection object as a parameter. This Connection object represents a valid connection to a RabbitMQ server, and is used to create new Channel objects.

RMQChannelPool:

The RMQChannelPool class is designed to be used by multiple threads simultaneously, so it uses a BlockingQueue to ensure thread-safety. The borrowObject() method blocks until a channel is available, while the returnObject() method adds the channel back to the pool for other threads to use. The purpose of this class is to provide a simple way to manage a fixed number of RabbitMQ channels in a multi-threaded environment, which can improve performance and scalability of applications that use RabbitMQ.

SwipeServlet:

In my SwipeServlet class, the new feature I added is that the SwiperServlet class determines whether the swipe was a "like" or "dislike" and sends the swipe data to a RabbitMQ message queue using the RabbitMQ Java client library. In the constructor, a connection is established with a RabbitMQ server using the "ConnectionFactory" class. The RabbitMQ server's IP address, virtual host, and login credentials are set in the ConnectionFactory object. A new connection is established using the ConnectionFactory object, and a new channel factory is created using the Connection object. The RMQChannelPool object is then instantiated using the channel factory object and a POOL_SIZE of 20.

The **sendMsgToQueue()** method sends the swipe data to the RabbitMQ queue. It retrieves a channel from the `RMQChannelPool` and declares an exchange named "swipe_exchange" using the channel. It then serializes the `SwipeDetails` object to JSON and publishes it to the exchange using the channel. And this method is called after valid a body based on if swipes right or left which represents true or false respectively.

In my consumer side, a connection to the RabbitMQ server is established using the `Connection` class from the RabbitMQ client library. A channel is then created using the `createChannel()` method, which is used to interact with the RabbitMQ server. The `queueDeclare()` method is used to declare a queue for the consumer to consume messages from, and the `queueBind()` method is used to bind the queue to an exchange so that messages can be routed to the queue.

The `basicConsume()` method is used to start consuming messages from the queue. When a message is received, the `DeliverCallback` that was passed to the `basicConsume()` method is invoked with the delivery details and the message body. And then the message is deserialized from JSON to a `SwipeDetails` object using the `Gson` library and the `basicAck()` method is then called to acknowledge receipt of the message and indicating that can be removed it from the queue.

3. Test run results (command lines showing metrics, RMQ management windows showing queue size, send/receive rates) for a single servlet showing your best throughput.

Throughput:

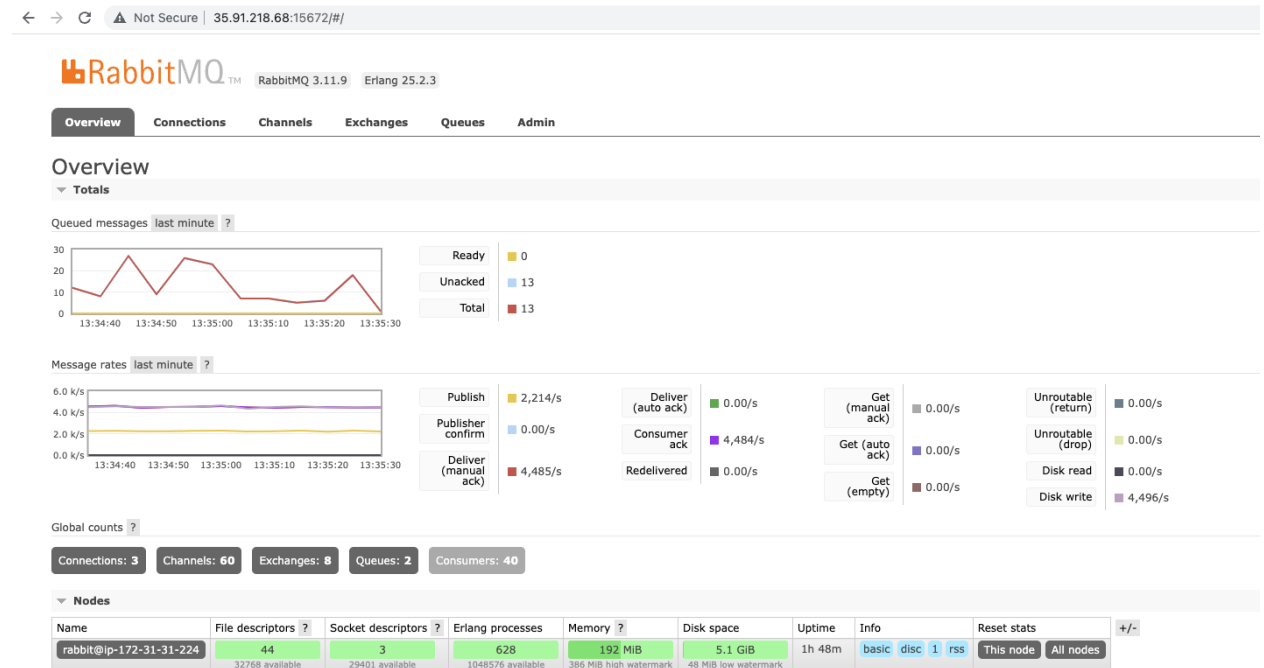
A screenshot of a Java application output window titled "MultithreadClient". The window shows the following output:

```
/Library/Java/JavaVirtualMachines/jdk-11.0.11.jdk/Contents/Home/bin/java ...  
Number of successful requests: 500000  
Number of fail requests: 0  
walltime: 231 seconds  
Total throughput: 2164 req/s  
  
Process finished with exit code 0
```

RMQ queue size, send/receive rates:

Queue 1: about 2271/s (deliver) , 2263/s(incoming)

Queue2: about 2271/s (deliver) , 2263/s(incoming)



← → ↻ ⚠ Not Secure | 35.91.218.68:15672/#/queues

RabbitMQ™

RabbitMQ 3.11.9 Erlang 25.2.3

Overview Connections Channels Exchanges Queues Admin

Queues

▼ All queues (2)

Pagination

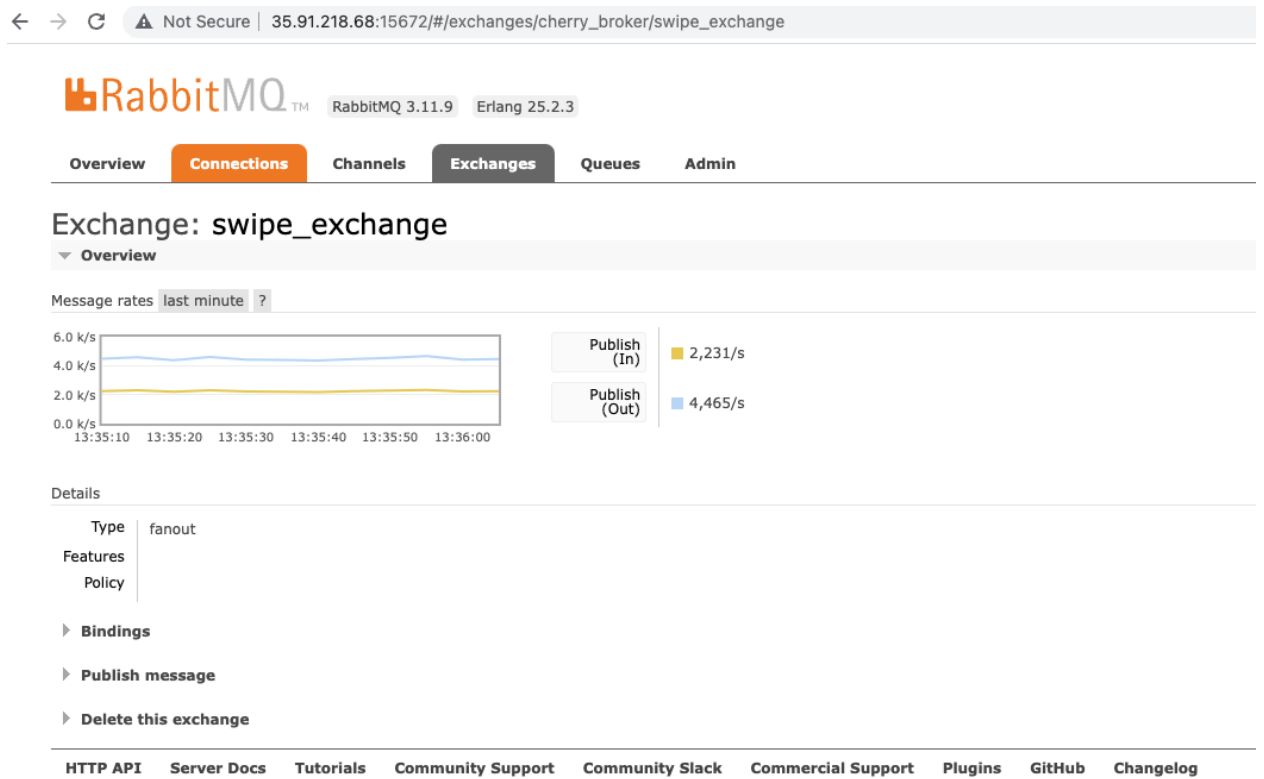
Page 1 of 1 - Filter: ☐ Regex ?

Overview				Messages			Message rates				+/-
Name	Type	Features	State	Ready	Unacked	Total	incoming	deliver / get	ack		
swipe_queue_1	classic	D	running	0	18	18	2,263/s	2,271/s	2,270/s		
swipe_queue_2	classic	D	running	0	16	16	2,263/s	2,271/s	2,270/s		

► Add a new queue

HTTP API Server Docs Tutorials Community Support Community Slack Commercial Support Plugins GitHub Changelog

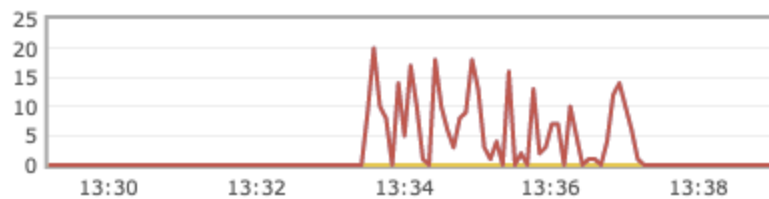
Exchange rate:



Queue swipe_queue_1

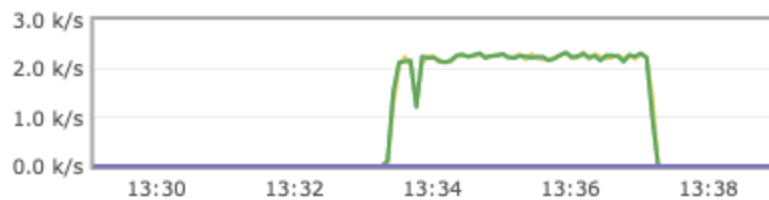
▼ **Overview**

Queued messages last ten minutes ?



Ready	
Unacked	
Total	

Message rates last ten minutes ?



Publish	
Deliver (manual ack)	
Deliver (auto ack)	

Details



RabbitMQ 3.11.9

Erlang 25.2.3

Overview

Connections

Channels

Exchanges

Queues

Queue swipe_queue_2

▼ Overview

Queued messages last ten minutes ?

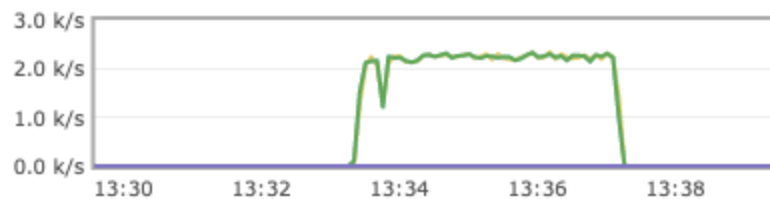


Ready

Unacked

Total

Message rates last ten minutes ?



Publish

Deliver
(manual
ack)

Deliver
(auto ack)

Details

Features | durable: true

State | idle

4. Test run results (command lines showing metrics, RMQ management windows showing queue size, send/receive rates) for a load balanced servlet showing your best throughput.

```
src
├── main
│   └── java
│       └── org.example
└── ...

MultiThreadClient.java
1 public class MultiThreadClient {
2     private static final int numOfThreads = 110;
3     private static AtomicInteger numSuccessfulRequests = new AtomicInteger(0);
4
5     public static void main(String[] args) {
6         // ...
7     }
8 }
```

MultiThreadClient

/Library/Java/JavaVirtualMachines/jdk-11.0.11.jdk/Contents/Home/bin/java ...

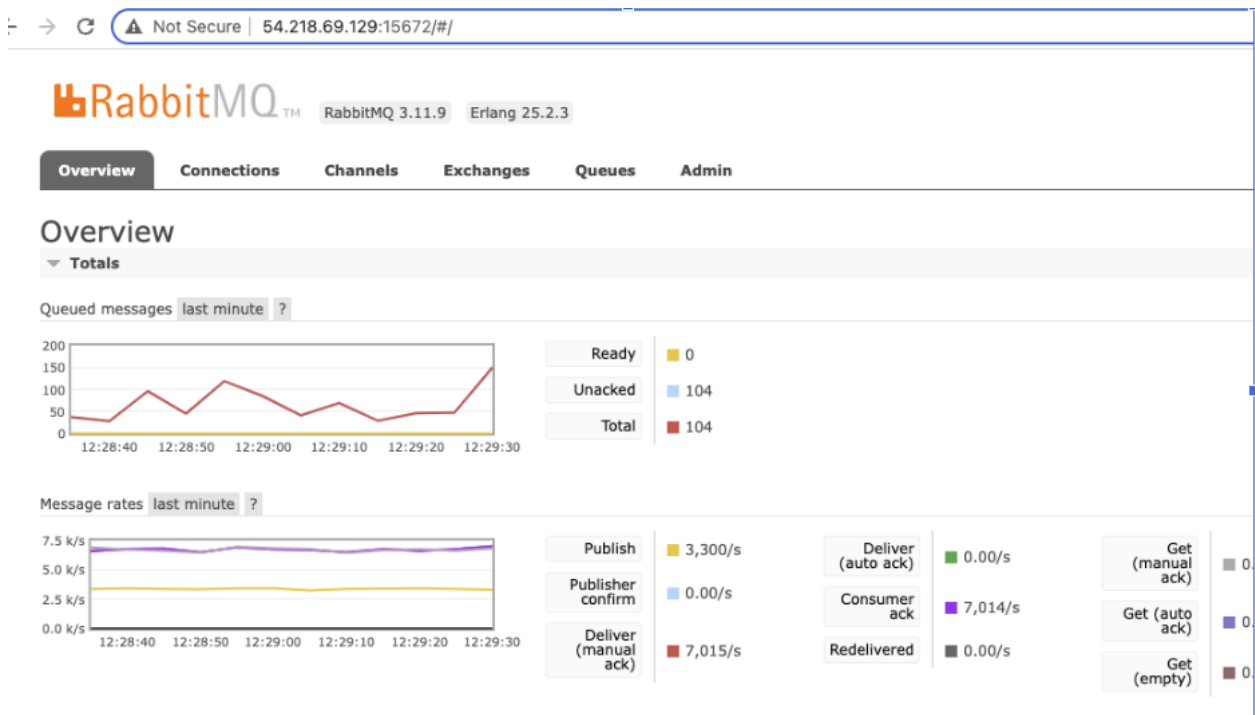
Number of successful requests: 500000

Number of fail requests: 0

walltime: 158 seconds

Total throughput: 3164 req/s

Process finished with exit code 0




Queue swipe_queue_1

▼ Overview

Queued messages

last minute

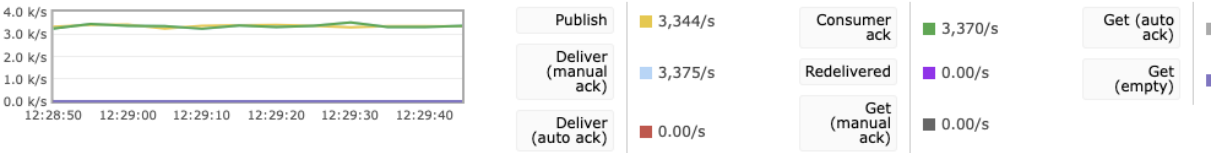
?



Message rates

last minute

?

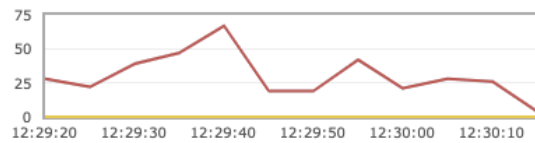


Details									
Features	durable: true	State	<div><div></div> running</div>			Total	Ready	Unacked	In
	Policy	Consumers	100	Messages ?		27	0	27	
	Operator policy	Consumer capacity ?	100%	Message body bytes ?		2.6 KiB	0 B	2.6 KiB	
	Effective policy definition			Process memory ?		4.6 MiB			

Queue swipe_queue_2

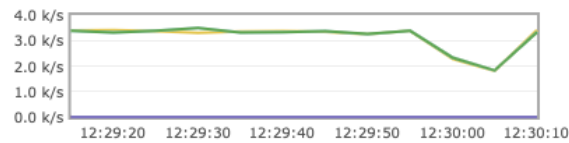
▼ Overview

Queued messages last minute ?



Ready 0
Unacked 4
Total 4

Message rates last minute ?



Publish 3,412/s
Deliver (manual ack) 3,335/s
Deliver (auto ack) 0.00/s

Consumer ack 3,324/s
Redelivered 0.00/s
Get (manual ack) 0.00/s

Details

Features	durable: true	State	running	Messages ?	Total	Read
Policy		Consumers	100		4	1
Operator policy		Consumer capacity ?	100%	Message body bytes ?	410 B	0 B

Overview

▼ Totals

Queued messages last ten minutes ?


Ready

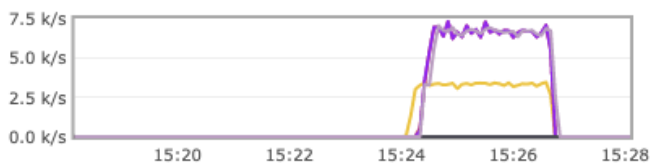
0

Unacked

0

Total

0

Message rates last ten minutes ?


Publish

0.00/s

Publisher confirm

0.00/s

Deliver (manual ack)

0.00/s

Deliver (auto ack)

Consumer ack

Redelivered

Global counts ?

Queue swipe_queue_1

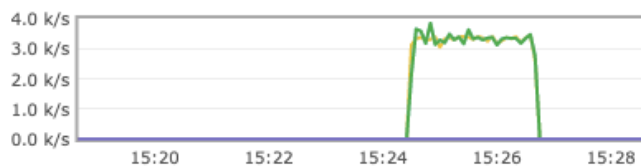
▼ Overview

Queued messages [last ten minutes](#) [?](#)



Ready	0
Unacked	0
Total	0

Message rates [last ten minutes](#) [?](#)



Publish	0.00/s
Deliver (manual ack)	0.00/s
Deliver (auto ack)	0.00/s

Consumer ack
Redelivered
Get (manual ack)

Details

Queue swipe_queue_2

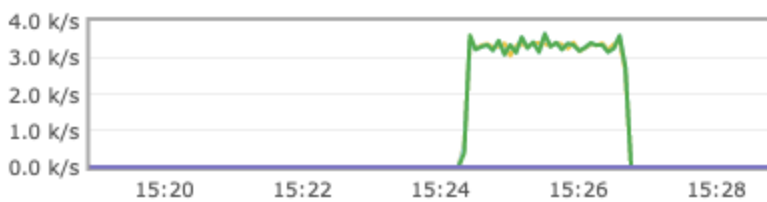
Overview

Queued messages **last ten minutes** ?



Ready	0
Unacked	0
Total	0

Message rates **last ten minutes** ?



Publish	0.00/s
Deliver (manual ack)	0.00/s
Deliver (auto ack)	0.00/s

Details

Features	durable: <u>true</u>	State	<u>idle</u>	
Policy		Consumers	20	
Operator policy		Consumer capacity ?	100%	
Effective policy definition				

The results clearly show that incorporating an Application Load Balancer has increased the throughput by approximately 50% compared to using a single Tomcat server. Additionally, the producer and consumer rates also increased by about 50% when compared to using a single servlet. The queue size remained consistently low in both the single servlet and ALB modes.

It is essential to recognize that including an Application Load Balancer did not double the throughput, since various factors can affect performance aside from just adding a server. Moreover, even after sending 500,000 requests in single servlet mode, the CPU usage remained consistently low, around 20-30%. Notably, the CPU usage for each servlet instance also remained similar when operating under the ALB mode.

Tune the system & Loading test

How many client threads are optimal to maximize system throughput?

To optimize system throughput, I conducted various tests on the number of client threads and queue consumer threads. Based on my observations, I found that around 110 client threads work best for maximizing throughput. Additionally, to keep the queue size as low as possible, I found that having 20 queue consumer threads was optimal. And the channel pool size for my producer is 30 working best.

How many queue consumers threads are needed to keep the queue size as close to zero as possible?

During the tuning process, I made adjustments to various configuration parameters such as the number of threads in the client, the channel pool size for producers, and the number of consumer threads. To accurately measure the performance, I performed tests at different times to account for external factors like AWS network fluctuation.

After experimenting with different configurations, I found the best throughput, producer, and consumer rates by adjusting one value at a time and monitoring the results. Overall, my testing helped me to identify the appropriate number of client threads, producer channel pool size, and consumer threads for optimal system performance.